

YOLO

"You Only Look Once" –
Object Detection



Group 6: Joaquin Arias, Era Gërbeshi,
Mohamed Mourad Ouazghire, Naomi Mukuhi



What is our project?

Real-Time Object Detection

Locating and identifying:

- Duckiebots
- Ducks
- Stop signs
- Road signs

Real-world application:

- Autonomous driving



DUCKIETOWN

X

YOLO

Object Detection

What is it?

Technique to identify and locate objects in an image or video.

Applications

- Surveillance
- Self-driving cars
- Robotics






Image Classification vs. Object Detection

Classification



Cat

Detection



Cat, Dog, Dog

V7 Labs

Object Detection

Two-Stage

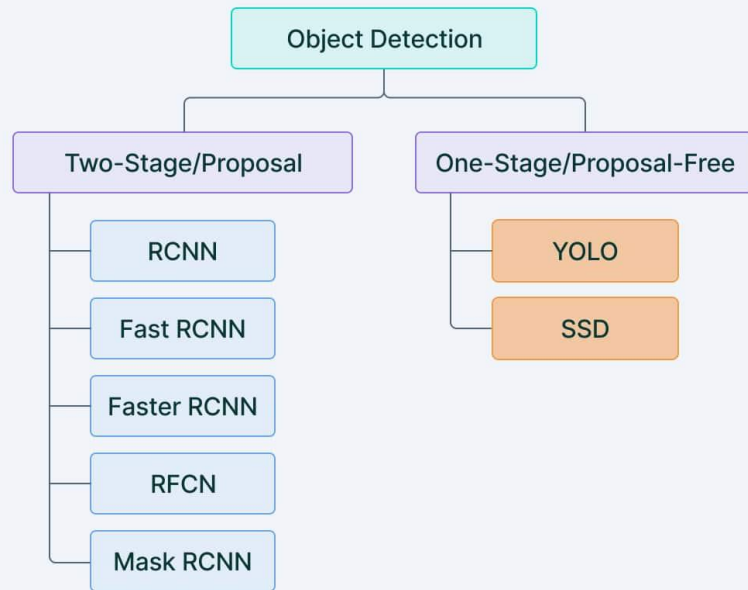
- Two passes of the input image:
 - 1st: Proposal of set of regions that might contain the object.
 - 2nd: classify and refine regions.
- Computationally expensive.

One-Stage

- One pass of the input image.
- Predicts the bounding step without region proposal.
- Computationally efficient and faster.



One and two stage detectors



You Only Look Once




- Widely adopted.
- Easy to implement and adapt to new tasks due to its simple architecture.
- Struggles with small objects

Fast

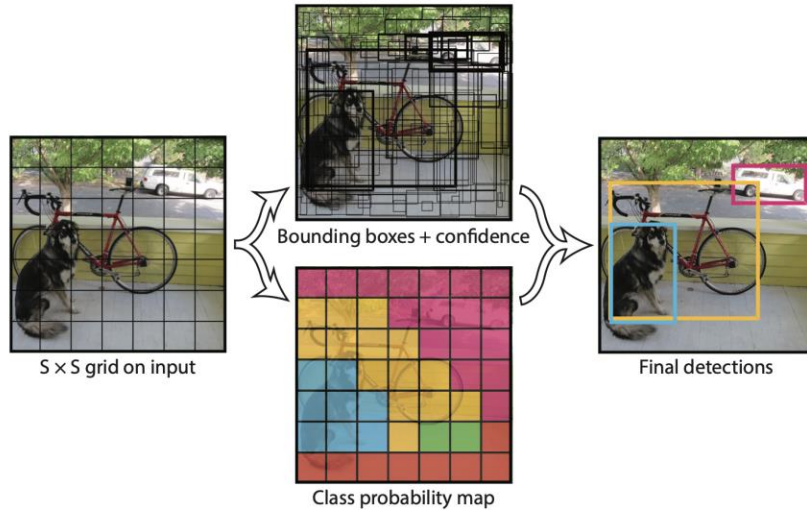
Efficient

Simple



(Hui, 2018)

How does it work?




(Redmon et al., 2016)

The Model

- Divide image into an S x S grid.
- Predict B bounding boxes and the confidence for them.
- Predict class probabilities.
- Output final detections.

The Architecture



(Redmon et al., 2016)

Convolutional Layers

Extract features and create features map.

24

Connected Layers


Predict bounding boxes and class probabilities for each cell in the grid.

2

Output

A list of bounding boxes with their respective probabilities.

1



Implementation



DUCKIETOWN

X

YOLO

Implementation

+500 similar images




The Dataset

- 507 compressed images from the duckiebot's camera, capturing the different road signs, duckies and duckiebots during its movement.


Separation of blurry images from non-blurry ones.

The Dataset

- For optimizing our object detection efficiency, we built a python script to only conserve the non-blurry images, increasing the model training quality, and reducing the training time span.



- With a filtering threshold of 200, the total number of images from the dataset gets reduced from +1500->507



Kept



Removed



The Laplacian Filter was used for this purpose

detect_blurry_img.py

```
import argparse
from pathlib import Path
import sys

def main():

    parser = argparse.ArgumentParser(description="Eliminate blurry pictures")
    parser.add_argument("inputFolder", help="Path of folder containing images to classify", type=str)
    parser.add_argument("blurryFolder", help="Path of folder where blurry images will be sent", type=str)
    parser.add_argument("notBlurryFolder", help="Path of folder where non blurry images will be sent", type=str)
    parser.add_argument("--threshold", help="Threshold for blurry detection, default is 200", type=int, default=200)

    args = parser.parse_args()

    data_folder = Path(args.inputFolder)
    blurry_folder = Path(args.blurryFolder)
    good_folder = Path(args.notBlurryFolder)
    threshold = args.threshold

    if not data_folder.is_dir():
        print("{} is not a directory".format(data_folder))
        sys.exit(1)

    if not blurry_folder.is_dir():
        print("{} is not a directory".format(blurry_folder))
        sys.exit(1)

    if not good_folder.is_dir():
        print("{} is not a directory".format(good_folder))
        sys.exit(1)

    # Recognize jpg or jpeg images
    images = list(data_folder.glob('*.jpg'))
    images.extend(list(data_folder.glob('*.jpeg')))

    # Go through all images in data folder
    for imageFile in images:
        print('Processing image {}'.format(imageFile))

        image = cv.imread(str(imageFile))
        gray = cv.cvtColor(image, cv.COLOR_BGR2GRAY)
        fm = cv.Laplacian(gray, cv.CV_64F).var()

        # blurry
        if fm < threshold:
            cv.imwrite(str(blurry_folder.joinpath(imageFile.name)), image)
        # not blurry
        else:
            cv.imwrite(str(good_folder.joinpath(imageFile.name)), image)

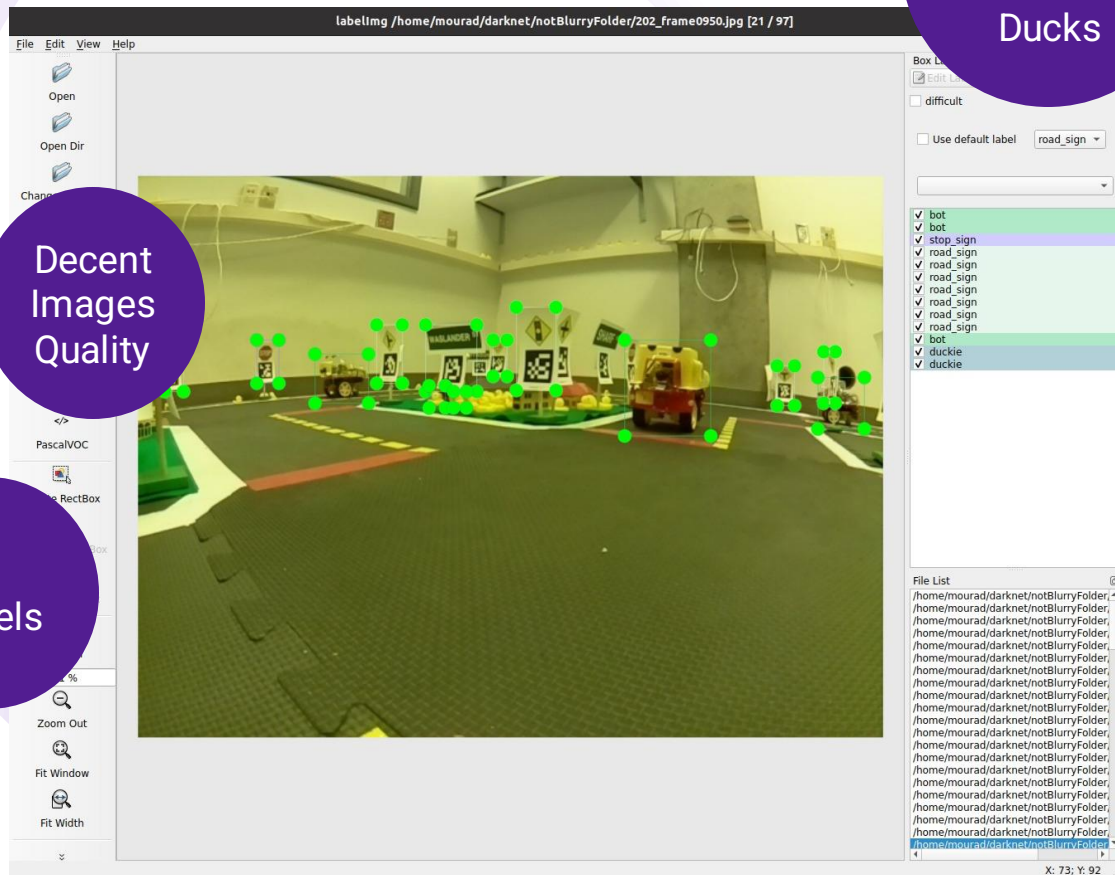
if __name__ == "__main__":
    main()
```

Labeling

Stop Signs
Road Signs
Duckiebots
Ducks

Decent
Images
Quality

4
Labels




```
Files + Code + Text Copy to Drive RAM Disk
[x] darknet
sample_data
cuda_9.2.88_396.26_linux.run

488255: 0.129388, 1.103002 avg, 0.000010 rate, 0.010538 seconds, 488255 images
Loaded: 0.007381 seconds
Region 16 Avg IOU: -nan, Class: -nan, Obj: -nan, No Obj: 0.000000, .5R: -nan, .75R: -nan, count: 0
Region 23 Avg IOU: 0.741594, Class: 0.999867, Obj: 0.611886, No Obj: 0.004983, .5R: 1.000000, .75R: 0.636364, count: 11
488256: 3.663269, 1.359029 avg, 0.000010 rate, 0.010714 seconds, 488256 images
Loaded: 0.006403 seconds
Region 16 Avg IOU: 0.862634, Class: 0.999968, Obj: 0.999676, No Obj: 0.001831, .5R: 1.000000, .75R: 1.000000, count: 1
Region 23 Avg IOU: 0.875104, Class: 0.999929, Obj: 0.770646, No Obj: 0.004302, .5R: 1.000000, .75R: 1.000000, count: 6
488257: 1.183214, 1.341447 avg, 0.000010 rate, 0.011311 seconds, 488257 images
Loaded: 0.006265 seconds
Region 16 Avg IOU: 0.829091, Class: 0.998652, Obj: 0.097909, No Obj: 0.000276, .5R: 1.000000, .75R: 1.000000, count: 1
Region 23 Avg IOU: 0.881492, Class: 0.999885, Obj: 0.892967, No Obj: 0.001632, .5R: 1.000000, .75R: 1.000000, count: 3
488258: 0.565657, 1.263868 avg, 0.000010 rate, 0.022542 seconds, 488258 images
Loaded: 0.001222 seconds
Region 16 Avg IOU: -nan, Class: -nan, Obj: -nan, No Obj: 0.000000, .5R: -nan, .75R: -nan, count: 0
Region 23 Avg IOU: 0.667293, Class: 0.999858, Obj: 0.606192, No Obj: 0.003199, .5R: 1.000000, .75R: 0.142857, count: 7
488259: 2.496946, 1.397176 avg, 0.000010 rate, 0.014784 seconds, 488259 images
Loaded: 0.004051 seconds
Region 16 Avg IOU: -nan, Class: -nan, Obj: -nan, No Obj: 0.000002, .5R: -nan, .75R: -nan, count: 0
Region 23 Avg IOU: 0.797511, Class: 0.957471, Obj: 0.797786, No Obj: 0.004292, .5R: 1.000000, .75R: 0.750000, count: 9
488260: 2.013806, 1.449839 avg, 0.000010 rate, 0.013265 seconds, 488260 images
Resizing
608
Loaded: 0.000054 seconds
Region 16 Avg IOU: 0.819272, Class: 0.996566, Obj: 0.957306, No Obj: 0.004415, .5R: 1.000000, .75R: 1.000000, count: 13
Region 23 Avg IOU: 0.730033, Class: 0.999979, Obj: 0.928511, No Obj: 0.001727, .5R: 1.000000, .75R: 0.923077, count: 13
488261: 0.939160, 1.398771 avg, 0.000010 rate, 0.015260 seconds, 488261 images
Loaded: 0.006780 seconds
Region 16 Avg IOU: -nan, Class: -nan, Obj: -nan, No Obj: 0.000000, .5R: -nan, .75R: -nan, count: 0
Region 23 Avg IOU: 0.680077, Class: 0.973570, Obj: 0.710726, No Obj: 0.003078, .5R: 0.923077, .75R: 0.923077, count: 13
488262: 3.482118, 1.607106 avg, 0.000010 rate, 0.021347 seconds, 488262 images
Loaded: 0.003632 seconds
Region 16 Avg IOU: -nan, Class: -nan, Obj: -nan, No Obj: 0.000000, .5R: -nan, .75R: -nan, count: 0
Region 23 Avg IOU: 0.806131, Class: 0.999795, Obj: 0.683546, No Obj: 0.002504, .5R: 1.000000, .75R: 0.750000, count: 9
488263: 2.347107, 1.681106 avg, 0.000010 rate, 0.025316 seconds, 488263 images
Loaded: 0.004836 seconds
```

Training

- Prepare the data for training:
 - Training, Testing, Evaluation sets.
- Prepare config. Files:
 - Classes, model selection, determine classes.
- Fork Darknet repository



Ground Truth Bounding Boxes

| | | | | |
|---|------|------|------|------|
| 3 | 0.19 | 0.22 | 0.04 | 0.13 |
| 3 | 0.23 | 0.21 | 0.03 | 0.13 |
| 3 | 0.62 | 0.2 | 0.05 | 0.11 |
| 3 | 0.7 | 0.21 | 0.06 | 0.15 |
| 2 | 0.96 | 0.28 | 0.03 | 0.12 |
| 1 | 0.34 | 0.27 | 0.03 | 0.04 |

Label

```
datasets/testset/frame_002166_png.rf.a3884e8d84306dec7b86dba2ce4ebca7.jpg  
road_sign:Left=319, Top=130, Right=341, Bottom=96
```

```
duckie:Left=243, Top=214, Right=283, Bottom=180
```

```
road_sign:Left=154, Top=144, Right=172, Bottom=105
```

```
road_sign:Left=26, Top=196, Right=84, Bottom=169
```

```
datasets/testset/202_frame0806.jpg  
stop_sign:Left=86, Top=196, Right=118, Bottom=120
```

```
road_sign:Left=115, Top=192, Right=140, Bottom=115
```

```
bot:Left=182, Top=208, Right=240, Bottom=156
```

```
duckie:Left=419, Top=182, Right=451, Bottom=153
```

```
road_sign:Left=371, Top=192, Right=409, Bottom=124
```

```
bot:Left=425, Top=244, Right=521, Bottom=148
```

```
bot:Left=528, Top=240, Right=598, Bottom=163
```

```
duckie:Left=608, Top=216, Right=620, Bottom=196
```

Editing the source code

```
layer    filters  size      input           output
  0 conv     32  3 x 3 / 1   416 x 416 x  3   ->  416 x 416 x  32  0.299 BFLOPs
  1 conv     64  3 x 3 / 2   416 x 416 x 32   ->  208 x 208 x  64  1.595 BFLOPs
  .....
105 conv    255  1 x 1 / 1    52 x  52 x 256   ->   52 x  52 x 255  0.353 BFLOPs
106 detection


truth_thresh: Using default '1.000000'
Loading weights from yolov3.weights...Done!
data/dog.jpg: Predicted in 0.029329 seconds.
dog: 99%
truck: 93%
bicycle: 99%
```

Before

```
datasets/testset/204_frame0056.jpg: Predicted in 0.179878 seconds.
stop_sign: 88%
Bounding Box: Left=196, Top=112, Right=237, Bottom=190
stop_sign: 70%
Bounding Box: Left=167, Top=148, Right=181, Bottom=175
duckie: 95%
Bounding Box: Left=145, Top=197, Right=166, Bottom=216
duckie: 94%
Bounding Box: Left=89, Top=192, Right=112, Bottom=216
duckie: 87%
Bounding Box: Left=163, Top=196, Right=186, Bottom=219
duckie: 65%
Bounding Box: Left=473, Top=196, Right=491, Bottom=215
bot: 98%
Bounding Box: Left=0, Top=147, Right=69, Bottom=243
```

After

Results



Class: bot
True Positives (TP): 0
True Negatives (TN): 0
False Positives (FP): 0
False Negatives (FN): 44


Class: duckie
True Positives (TP): 22
True Negatives (TN): 0
False Positives (FP): 6
False Negatives (FN): 68

Class: road_sign
True Positives (TP): 84
True Negatives (TN): 0
False Positives (FP): 15
False Negatives (FN): 40

Class: stop_sign
True Positives (TP): 30
True Negatives (TN): 0
False Positives (FP): 1
False Negatives (FN): 13



Intersection



$$\text{Base} = \min(R_1, R_2) - \max(L_1, L_2)$$

$$\text{Height} = \min(T_1, T_2) - \max(B_1, B_2)$$

$$\text{Area} = \text{Base} \times \text{Height}$$

$$\text{IOU} = \frac{\text{Area}}{\text{Area}}$$

Union:

$$\text{Area} = (T_1 - B_1)(R_1 - L_1) + (T_2 - B_2)(R_2 - L_2) - \text{Area}$$

```
intersection_x1 = max(x1, x3)
intersection_y1 = min(y1, y3)
intersection_x2 = min(x2, x4)
intersection_y2 = max(y2, y4)

intersection_area = max(0, intersection_x2 - intersection_x1) * max(0, intersection_y1 - intersection_y2)
#print(f"({intersection_x2}-{intersection_x1}) * ({intersection_y1}-{intersection_y2})={intersection_area}")

box1_area = abs((x2 - x1) * (y2 - y1))
box2_area = abs((x4 - x3) * (y4 - y3))

union_area = box1_area + box2_area - intersection_area
#print(f"area={box1_area}+{box2_area}={union_area}")

iou = intersection_area / union_area
#print(f"iou={intersection_area}/{union_area}={iou}")
```

Results




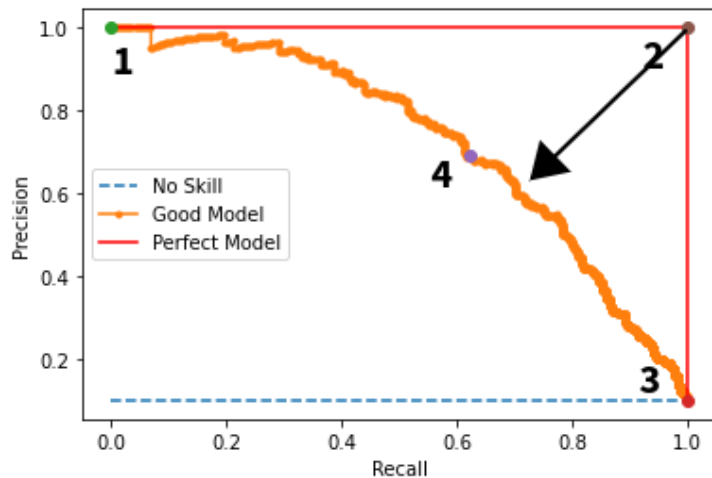
Precision: how often does the model predict correctly?

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

Recall: has the model predicted every time that it should have predicted?

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

Average precision: Area under the curve.



Results

Reference


Results

Average precision

| Class | Approx. Area Under the Curve (AP) |
|-----------|-----------------------------------|
| bot | 0.00 |
| duckie | 0.77 |
| road_sign | 0.65 |
| stop_sign | 0.75 |

Mean Average Precision (mAP) = 0.54

$$mAP = \frac{1}{N} \sum_{i=1}^N AP_i$$



Challenges and Recommendations

Assessment Problems

- The "Bot" class only showed False Negatives (FN), and the model failed to detect any instances of this class

Duckiebot

- Wiring problems caused problems turning it on and off.
- Problems with Jetson Nano OS

Work process

Terminal 1

```
export ROS_MASTER_URI=http://[name].local:11311

source ~/svo_ws/devel/setup.bash

roslaunch image_transport republish compressed
in:=/[name]/camera_node/image raw
out:=/[name]/camera_node/image/raw
```

darknet_ros node

- Problems accessing the node



THANK YOU

*"You only live once - YOLO."
:)*

